

Master's Seminar: Artificial Intelligence - Visual Computing

3D Gaussian Splatting

Sebastian Frehmel

Artificial Intelligence Group,
Fernuniversität Hagen, Germany

September 2nd, 2026



► The NeRF Dilemma

Neural Radiance Fields (NeRFs) are either slow and provide photorealistic scenes or they are reasonably fast but have to sacrifice image quality. Real-time applications are not possible.

► 3D Gaussians As The Next Logical Step

Using 3D Gaussians as the explicit base primitive, cleverly decomposing them to make fast backpropagation possible, and abandoning neural networks is the key to unlocking previously unseen rendering speeds.

► Engineering Effort

State-of-the-art GPU integration yields massive speed-ups, mostly due to the explicit nature of the base primitive.

- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure
 - Point Cloud Initialization
 - Forward Pass
 - Backward Pass
- 4 Novel-View Synthesis
- 5 Summary and Discussion

Core Concepts

3D Scene Creation

Novel-View Synthesis

Splatting

Definitions

Creates 3D environments such as individual objects or outdoor spaces from an overlapping set of 2D input images.

Core Concepts

3D Scene Creation

Novel-View Synthesis

Splatting

Definitions

Synthesizes new 2D imagery from the 3D scene using new camera positions.

Real-time capability requires at least 24–30 frames per second (FPS) to appear smooth enough for the human eye.

Core Concepts


3D Scene Creation

Novel-View Synthesis


Splatting


Definitions


The process of moving all elements inside the view frustum onto the view plane preserving order, adding to aggregated color and opacity of each pixel.





- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure
 - Point Cloud Initialization
 - Forward Pass
 - Backward Pass
- 4 Novel-View Synthesis
- 5 Summary and Discussion













Multivariate Normal Distribution: $\alpha(p) = o \cdot \exp\left(-\frac{1}{2}(p - \mu)^T \Sigma^{-1}(p - \mu)\right)$



$$\text{Multivariate Normal Distribution: } \alpha(p) = o \cdot \exp\left(-\frac{1}{2}(p - \mu)^T \Sigma^{-1}(p - \mu)\right)$$



Multivariate Normal Distribution: $\alpha(p) = o \cdot \exp\left(-\frac{1}{2}(p - \mu)^T \Sigma^{-1}(p - \mu)\right)$

Mathematical Setup

The Validity Problem

Core Decomposition

Decomposed Components

Guaranteed Validity

Details

To work with gradient descent, the covariance matrix Σ must be differentiable and remain positive semi-definite (eigenvalues ≥ 0) to be geometrically valid.

Running a loop during backpropagation to find valid values is computationally too costly.

Mathematical Setup

The Validity Problem

Core Decomposition

Decomposed Components

Guaranteed Validity

Details

Kerbl et al. bypassed this by decomposing Σ into a scaling matrix S and rotational matrix R mirroring standard eigendecomposition ($\Sigma = Q\Delta Q^T$):

$$\Sigma = RSS^T R^T \text{ with } \Delta = S^2, Q = R$$

Mathematical Setup

The Validity Problem

Core Decomposition

Decomposed Components

Guaranteed Validity

Details

The diagonal scaling matrix S is further decomposed into three distinct scaling values. These represent the roots of the eigenvalues of Σ .

The rotational matrix R is decomposed into a four-value *quaternion*.

Mathematical Setup

The Validity Problem

Core Decomposition

Decomposed Components

Guaranteed Validity


Details

3DGS only stores and optimizes these 7 specific values for each 3D Gaussian covariance matrix.

Because an eigendecomposition is always positive semi-definite, constructing a new covariance matrix from these components is guaranteed to be geometrically valid.

- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure**
 - Point Cloud Initialization
 - Forward Pass
 - Backward Pass
- 4 Novel-View Synthesis
- 5 Summary and Discussion

The Training Cycle



- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure**
 - Point Cloud Initialization
 - Forward Pass
 - Backward Pass
- 4 Novel-View Synthesis
- 5 Summary and Discussion

Initialization Steps

Structure from Motion

Camera Matrix

Primitive Initialization

Details

The 3DGS algorithm requires initial geometry, which is commonly derived from a set of input images.

Because camera poses are usually unknown, Structure from Motion (SfM) software like COLMAP is used to solve a geometric optimization problem, providing an accurate 3D point cloud.

Initialization Steps

Structure from Motion

Camera Matrix

Primitive Initialization

Details

As another result from SfM, each input image is assigned a camera matrix W which contains the camera's orientation in world space.

Initialization Steps

Structure from Motion

Camera Matrix

Primitive Initialization


Details

The SfM points act as the origins for the initial 3D Gaussians.

- ▶ They are initialized as perfect spheres.
- ▶ Radii are deduced from their three nearest neighbors.
- ▶ Opacity is set to a low value, and color to any RGB value.

- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure**
 - Point Cloud Initialization
 - **Forward Pass**
 - Backward Pass
- 4 Novel-View Synthesis
- 5 Summary and Discussion

Forward Pass: Pipeline Architecture



Phase 1: Preparing Constituent Attributes

- ▶ The GPU is running one GPU thread per 3D Gaussian.
- ▶ Raw values optimized during backpropagation are converted.
- ▶ Opacity is bounded to $(0, 1)$ using a sigmoid function, scale uses the exponential function, and the rotational quaternion is normalized.

Phase 1: Preparing Constituent Attributes

- ▶ The GPU is running one GPU thread per 3D Gaussian.
- ▶ Raw values optimized during backpropagation are converted.
- ▶ Opacity is bounded to $(0, 1)$ using a sigmoid function, scale uses the exponential function, and the rotational quaternion is normalized.

Phase 2: Reconstructing the 3D Covariance Matrix


- ▶ The algorithm mathematically rebuilds the 3D geometric shape.
- ▶ It combines the scaling matrix S and rotational matrix R to form the full 3D covariance matrix:

$$\Sigma_{3D} = RSS^T R^T$$


Phase 3: Projection from 3D to 2D

- ▶ Converts 3D Gaussians (ellipsoids) in world space into flat 2D Gaussians (ellipses) in image space.
- ▶ Uses a linear affine approximation to keep the 2D Gaussian elliptic (Jacobian J).
- ▶ The 3D matrix is rotated into view space via camera matrix W and then projected:

$$\Sigma_{2D} = JW\Sigma_{3D}W^TJ^T$$




Forward Pass: Pipeline Architecture



Phase 4: Tile Mapping (Preparing Radix Sort)

- ▶ The target output image is divided into a grid of tiles of 16×16 pixels.
- ▶ The algorithm calculates a 2D bounding box to map exactly which tiles a 2D Gaussian covers.
- ▶ Assigns each 2D Gaussian to a tile. If overlapping, assigns to all affected tiles.




Phase 4: Tile Mapping (Preparing Radix Sort)

- ▶ The target output image is divided into a grid of tiles of 16×16 pixels.
- ▶ The algorithm calculates a 2D bounding box to map exactly which tiles a 2D Gaussian covers.
- ▶ Assigns each 2D Gaussian to a tile. If overlapping, assigns to all affected tiles.

Phase 5: Global Radix Sort

- ▶ The GPU waits until all per-Gaussian computations are finished.
- ▶ A highly optimized Radix Sort is executed, organizing the 2D Gaussians for each tile strictly by depth (closest to farthest).



Phase 6: Tile-Based Rasterization (Splatting)


- ▶ Operating in 2D image space, the GPU assigns one processing thread per pixel.
- ▶ **Alpha Compositing:** Pixels accumulate color by blending the sorted list of overlapping 2D Gaussians front-to-back:

$$C_p = \sum_{i=1}^N c_i \alpha_i T_i, \quad \text{with} \quad T_i = \prod_{j=1}^{i-1} (1 - \alpha_j) \quad \text{the transmittance}$$


Phase 6: Tile-Based Rasterization (Splatting)

- ▶ Operating in 2D image space, the GPU assigns one processing thread per pixel.
- ▶ **Alpha Compositing:** Pixels accumulate color by blending the sorted list of overlapping 2D Gaussians front-to-back:

$$C_p = \sum_{i=1}^N c_i \alpha_i T_i, \quad \text{with} \quad T_i = \prod_{j=1}^{i-1} (1 - \alpha_j) \quad \text{the transmittance}$$




The Training Cycle




- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure**
 - Point Cloud Initialization
 - Forward Pass
 - **Backward Pass**
- 4 Novel-View Synthesis
- 5 Summary and Discussion


Backward Pass: Architecture & Gradient Flow



Backward Pass: Architecture & Gradient Flow




Backward Pass: Architecture & Gradient Flow




Adaptive Density Control

- ▶ Steps in every 100 iterations to counter over-reconstruction (covering too much geometry) and under-reconstruction (too few Gaussians).
- ▶ **Splitting:** Over-reconstructing Gaussians are split into two smaller ones.
- ▶ **Cloning:** Under-reconstructing Gaussians are cloned.

The Training Cycle



- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure
 - Point Cloud Initialization
 - Forward Pass
 - Backward Pass
- 4 Novel-View Synthesis**
- 5 Summary and Discussion



- ▶ Because the scene geometry is finalized, the backward pass, Adaptive Density Control, and other components are completely disabled.
- ▶ Freed from this overhead and running at the raw speed of the tile-based rasterizer, 3DGS quickly reaches real-time rendering levels.

- 1 Background
- 2 3D Gaussian Data Structure
- 3 Training Procedure
 - Point Cloud Initialization
 - Forward Pass
 - Backward Pass
- 4 Novel-View Synthesis
- 5 Summary and Discussion**

- ▶ 3D Gaussian Splatting is a learnable and explicit radiance field representation for 3D Reconstruction, abandoning neural networks.
- ▶ Kerbl et al. combined multiple existing approaches and technologies, added new ideas, and mapped them onto modern GPU hardware.
- ▶ We analyzed the training procedure showing in detail how Kerbl. et al made it possible to have real-time rendering of reconstructed 3D scenes for the first time.

Thank you for your attention. I am looking forward to taking your questions.

6 Videos

7 Backpropagation Calculations

8 Selected References

Training in Progress

<https://youtu.be/F9rgbvhSad8?t=550>



<https://youtu.be/yd40zUnr7eQ>



Real-Time Rendering

Example from the original paper's project website

<https://repo-sam.inria.fr/fungraph/3d-gaussian-splatting/content/videos/bicycle.mp4>



6 Videos

7 Backpropagation Calculations

8 Selected References

Phase 3: 2D Splat to 2D Geometry (Opacity & Position)

- ▶ We recall the simplified opacity distribution: $\alpha_i(p) = o \cdot e^E$, where E is the exponent derived from the pixel distance Δ and the inverse 2D covariance matrix.
- ▶ **Base Opacity (o):** The derivative scales directly with the distance from the center of the 2D Gaussian (e^E):

$$\frac{\partial L}{\partial o} = \frac{\partial L}{\partial \alpha_i} \cdot e^E$$

- ▶ **Position (μ_{2D}):** Let vector $v = \Sigma_{2D}^{-1} \Delta p$. This vector points directly from the 2D Gaussian's center to the pixel:

$$\frac{\partial L}{\partial \mu_{2D}} = \frac{\partial L}{\partial \alpha_i} \cdot \alpha_i(p) \cdot v$$

- ▶ If the evaluated opacity gradient is positive, the center is effectively pulled towards the pixel.

Phase 3: 2D Splat to 2D Geometry (Shape)

- ▶ **2D Covariance Matrix (Σ_{2D}):** We must calculate how the opacity at pixel p changes when the 2D shape of the Gaussian changes.
- ▶ This uses the same vector v from the position derivative, multiplied by its own transpose:

$$\frac{\partial L}{\partial \Sigma_{2D}} = \frac{\partial L}{\partial \alpha_i} \cdot \frac{1}{2} \alpha_i(p) \cdot (v \cdot v^T)$$

- ▶ Multiplying the vector by its transpose results in a matrix gradient.
- ▶ If the evaluated opacity gradient is positive, this matrix gradient causes the 2D Gaussian to stretch itself towards the pixel; otherwise, it causes it to contract.

Phase 4: 2D Geometry to 3D Geometry

- ▶ The GPU thread model changes to running per-Gaussian to unproject the 2D data back into 3D.
- ▶ This utilizes the multi-variable chain rule and the transposes of the forward projection matrices ($U = JW_{rot}$).
- ▶ **3D Position (μ_{3D}):** The gradient passes through view space (t) back to world space:

$$\frac{\partial L}{\partial \mu_{3D}} = W_{rot}^T J^T \frac{\partial L}{\partial \mu_{2D}} = U^T \left(\frac{\partial L}{\partial \mu_{2D}} \right)$$

- ▶ **3D Covariance (Σ_{3D}):** The final gradient unprojects cleanly using matrix U on both sides:

$$\frac{\partial L}{\partial \Sigma_{3D}} = U^T \left(\frac{\partial L}{\partial \Sigma_{2D}} \right) U$$

Phase 5: 3D Geometry to Matrix Components

- ▶ Gradients flow through the covariance matrix assembly ($\Sigma_{3D} = MM^T$ with $M = RS$) into scale (S) and rotation (R).
- ▶ First, the algorithm derives the intermediate matrix M :

$$\frac{\partial L}{\partial M} = 2 \left(\frac{\partial L}{\partial \Sigma_{3D}} \right) M$$

- ▶ Then, it splits this gradient into the separate components:

$$\frac{\partial L}{\partial S} = R^T \left(\frac{\partial L}{\partial M} \right) \quad \text{and} \quad \frac{\partial L}{\partial R} = \left(\frac{\partial L}{\partial M} \right) S$$

- ▶ The scale gradient ($\frac{\partial L}{\partial s}$) is extracted from the diagonal of $\frac{\partial L}{\partial S}$, while the quaternion gradient ($\frac{\partial L}{\partial q}$) aggregates the partial derivatives of R .

Phase 5: Backpropagating to Latent Attributes

- ▶ To close the loop, the gradients pass backwards through their respective activation functions to reach the raw, latent space variables.
- ▶ **Scale (Exponential):** Backpropagating through exp requires multiplying by s :

$$\frac{\partial L}{\partial s_{raw}} = \frac{\partial L}{\partial s} \cdot s$$

- ▶ **Rotation (Normalization):** Uses the Jacobian matrix ($J_{||}$) of the normalization function, with $J_{||} = \frac{I - qq^T}{||q_{raw}||}$:

$$\frac{\partial L}{\partial q_{raw}} = J_{||} \cdot \frac{\partial L}{\partial q}$$

- ▶ **Base Opacity (Sigmoid):** Backpropagating through the sigmoid function:

$$\frac{\partial L}{\partial o_{raw}} = \frac{\partial L}{\partial o} \cdot o(1 - o)$$

6 Videos

7 Backpropagation Calculations

8 Selected References

- [1] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkuehler, and George Drettakis.
3D Gaussian Splatting for Real-Time Radiance Field Rendering.
ACM Transactions on Graphics, 42(4):1–14, August 2023.
- [2] Christoph Lassner and Michael Zollhöfer.
Pulsar: Efficient Sphere-based Neural Rendering.
In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1440–1449, June 2021.
- [3] Shuzhe Wang, Vincent Leroy, Yohann Cabon, Boris Chidlovskii, and Jerome Revaud.
DUST3R: Geometric 3D Vision Made Easy.
In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 20697–20709, June 2024.
- [4] Lee Alan Westover.
Splatting: A Parallel, Feed-Forward Volume Rendering Algorithm.
PhD thesis, The University of North Carolina at Chapel Hill, 1991.

- [5] Matthias Zwicker, Hanspeter Pfister, Jeroen van Baar, and Markus Gross.
EWA Volume Splatting.
In *Proceedings Visualization, 2001, VIS '01*, pages 29–36, San Diego, CA, USA, 2001. IEEE.
- [6] Matthias Zwicker, Hanspeter Pfister, Jeroen van Baar, and Markus Gross.
Surface Splatting.
In *Proceedings of the 28th Annual Conference on Computer Graphics and Interactive Techniques, SIGGRAPH '01*, pages 371–378, New York, NY, USA, 2001. Association for Computing Machinery.